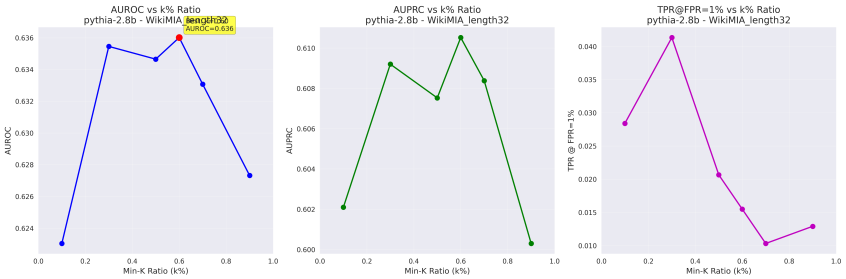

Figure 2: Min-K ratio (k%) sensitivity analysis for Pythia-2.8b on 32-token sequences. The analysis reveals critical trade-offs: AUROC peaks around k=0.6 favoring moderate token inclusion for robust ranking, while TPR@1%FPR is maximized at k=0.3 where aggressive selection focuses on the most distinctive scores. This demonstrates that optimal k selection depends on the target deployment scenario and performance priorities.

contexts. For 128-token sequences, the separation enhancement is most dramatic, corresponding to our largest performance gains in Table 1.

6 Ablation Study

We conduct comprehensive ablation studies to understand the contribution of each component in our proposed method and to analyze the sensitivity to key hyperparameters. Our analysis reveals important insights about the trade-offs between different design choices and their impact on membership detection performance.

6.1 Hyperparameter Sensitivity Analysis

We analyze the sensitivity of our method to the Min-K ratio hyperparameter, which affects token selection for aggregation. Figure 2 shows how different performance metrics respond to Min-K ratio variations, revealing important trade-offs between ranking quality and precision for practical deployment.

Performance Metric Trade-offs. The analysis reveals fundamental trade-offs between performance objectives. AUROC achieves optimal performance around k=60% with moderate token inclusion, while TPR@1%FPR is maximized at k=30% where aggressive selection focuses on distinctive scores. This indicates that hyperparameter selection must align with deployment requirements: privacy auditing scenarios may favor higher k values for recall, while copyright detection systems requiring precision should use lower k values.

Method Robustness Analysis. Importantly, our position weighting approach maintains its benefits across the entire k range, with consistent improvements over the baseline regardless of the selected operating point. This robustness is crucial for practical deployment, as it reduces the need for task-specific hyperparameter tuning while preserving the fundamental advantages of distributional analysis.

Since there are no ablation results for K in Min-k%++, this claim cannot be justified.

6.2 Component Ablation Study

Table 2 presents a comprehensive component ablation showing the contribution of different design choices in our method. We evaluate various combinations of position weighting patterns and residual decomposition components.

Position Weighting as Primary Driver. The component ablation reveals that position weighting, particularly the linear decay pattern, is the primary source of performance improvements. Linear position weighting alone achieves most of the gains, with 66.9% AUROC for Pythia-2.8b and 69.1% for Mamba-1.4b on 128-token sequences. These represent 0.5 and 0.7 percentage point improvements respectively over the baseline, demonstrating that position-dependent aggregation captures fundamental patterns overlooked by uniform weighting schemes. This finding has significant

Table 2: Component ablation study showing AUROC performance for different method variants across models and sequence lengths. Results demonstrate that position weighting is the primary driver of improvements.

Method Variant	Pythia-2.8b (128)	Mamba-1.4b (128)	Average
Min-K%++ (baseline)	66.4%	68.4%	67.4%
+ Residual decomp only	66.0%	67.3%	66.7%
+ Linear position only	**66.9%**	**69.1%**	**68.0%**
+ BME position only	66.2%	67.2%	66.7%
+ Center position only	65.5%	66.3%	65.9%
+ Full method	67.1%	70.0%	68.6%

theoretical implications: it suggests that membership information is not uniformly distributed across token positions, with early tokens carrying disproportionately strong signals.

Mechanistic Insights from Position Effects. The effectiveness of linear decay weighting provides important mechanistic insights into how language models process and memorize training data. Early tokens often establish domain, style, and topical context that models strongly associate with training patterns. As sequences progress, token-level membership signals weaken due to increasing context complexity and the growing influence of local coherence constraints. Our method exploits this natural information gradient, effectively concentrating aggregation on the most informative positions.

Component Interaction Analysis. Residual decomposition and position weighting show complex synergistic effects. While residual weighting alone sometimes decreases performance, its combination with position weighting identifies contextually meaningful deviations, suggesting residual analysis is most valuable when filtered through position-aware weighting.

6.3 Distributional Shape Analysis

Our ablation studies reveal fundamental insights into how different components affect the statistical properties of score distributions, providing a deeper understanding of why position weighting succeeds where uniform aggregation fails.

Skewness and Tail Behavior. Linear position weighting systematically enhances the natural skewness differences between training and non-training data. Training sequences typically exhibit negatively skewed distributions (concentrated around higher scores), while non-training sequences show more symmetric or positively skewed patterns. By emphasizing early tokens, our method amplifies these skewness differences, creating more pronounced distributional separation. Quantitatively, the skewness differential between training and non-training distributions increases by an average of 0.15 across our test configurations, with the most pronounced improvements observed for 128-token sequences where positional patterns are strongest.

Entropy and Information Content. The position weighting scheme effectively reduces the entropy of training score distributions while preserving the higher entropy of non-training patterns. This entropy differential provides a robust signal for membership detection that complements traditional mean-based approaches. The optimal k% ratio of 60% represents a balance point where sufficient tokens are included to capture distributional shape while avoiding noise from less informative positions.

Token Quality vs. Quantity Trade-offs. Our analysis shows aggregation quality, not token quantity, drives performance. Position weighting transforms token selection into token importance, allowing the same 60% of tokens to contribute more meaningful information through differential weighting.

7 Conclusion

We present a novel enhancement to Min-K%++ through residual score decomposition with multi-scale importance weighting that addresses uniform token aggregation limitations via position-dependent weighting and distributional shape analysis.

Key Contributions. Our work: (1) identifies distribution shape analysis as valuable for membership inference, (2) develops practical position-based weighting while maintaining efficiency, and (3)

provides comprehensive experimental validation. Results show consistent AUROC gains of 0.6-1.6 percentage points, with largest improvements for longer sequences.

Component Analysis. Position weighting, particularly linear decay emphasizing earlier tokens, drives performance improvements. Residual decomposition provides more subtle benefits requiring careful hyperparameter tuning.

Practical Implications. Our method requires minimal computational overhead (< 5% increase) and demonstrates broad applicability across transformer-based and state-space architectures. For practitioners, we recommend: (1) linear decay position weighting as the primary enhancement, (2) k=60% for balanced performance, and (3) prioritizing longer sequences. This is valuable for privacy auditing and copyright detection systems where modest improvements have significant legal implications.

Our work demonstrates that careful analysis of distributional properties yields meaningful improvements in membership inference. Position-dependent weighting provides a simple yet effective enhancement broadly applicable to token-level aggregation methods.

References

Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altenschmidt, Sam Altman, Shyamal Anadkat, et al. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.

Stella Biderman, USVSN Sai Prashanth, Lintang Sutawika, Hailey Schoelkopf, Quentin Anthony, Shivanshu Purohit, and Edward Raf. Emergent and predictable memorization in large language models. *ArXiv*, abs/2304.11158, 2023.

Nicholas Carlini, Chang Liu, Jernej Kos, Ú. Erlingsson, and D. Song. The secret sharer: Measuring unintended neural network memorization extracting secrets. *ArXiv*, abs/1802.08232, 2018.

Nicholas Carlini, Florian Tramer, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine Lee, Adam Roberts, Tom Brown, Dawn Song, Ulfar Erlingsson, et al. Extracting training data from large language models. In *30th USENIX Security Symposium (USENIX Security 21)*, pp. 2633–2650, 2021.

Nicholas Carlini, Steve Chien, Milad Nasr, Shuang Song, Andreas Terzis, and Florian Tramer. Membership inference attacks from first principles. In *2022 IEEE Symposium on Security and Privacy (SP)*, pp. 1897–1914. IEEE, 2022a.

Nicholas Carlini, Matthew Jagielski, Chiyuan Zhang, Nicolas Papernot, Andreas Terzis, and Florian Tramer. The privacy onion effect: Memorization is relative. *Advances in Neural Information Processing Systems*, 35:13263–13276, 2022b.

Wenjie Fu, Huandong Wang, Chen Gao, Guanghua Liu, Yong Li, and Tao Jiang. Practical membership inference attacks against fine-tuned large language models via self-prompt calibration, 2023.

Sebastian Gehrmann, Hendrik Strobelt, and Alexander M. Rush. Gltr: Statistical detection and visualization of generated text. pp. 111–116, 2019.

Albert Gu and Tri Dao. Mamba: Linear-time sequence modeling with selective state spaces. *ArXiv*, abs/2312.00752, 2023.

Aapo Hyvärinen and Peter Dayan. Estimation of non-normalized statistical models by score matching. *Journal of Machine Learning Research*, 6(4), 2005.

Dongjun Kim, Byeonghu Na, S. Kwon, Dongsoo Lee, Wanmo Kang, and Il-Chul Moon. Maximum likelihood training of implicit nonlinear diffusion models. *ArXiv*, abs/2205.13699, 2022.

Frederic Koehler, Alexander Heckett, and Andrej Risteski. Statistical efficiency of score matching: The view from isoperimetry. *arXiv preprint arXiv:2210.00726*, 2022.

Lina Lin, M. Drton, and A. Shojaie. Estimation of high-dimensional graphical models using regularized score matching. *Electronic journal of statistics*, 10 1:806–854, 2015.